

# DAT278 Lab 1 Manual

November 2023

Hari Abram, Mateo Vázquez Maceiras

## PREFACE

One major pillar of sustainable computing is maximizing energy efficiency, i.e., reducing the amount of energy consumed to carry out your prescribed work. To start your journey in sustainable computing, you will form groups of three, and explore multiple ways to improve the energy efficiency of running image recognition application. To join a group, go to Canvas > People > Lab Groups, and add yourself to one of the groups. If all the groups are taken, or you cannot find a group partner, please contact the Teaching Assistants (TAs) for help.

There are four different labs for you to do, and five sessions, as shown in Table I. Each lab is designed to take one session, but the fifth session can be used as extra time. If you have time enough, you can aim to get bonus points by doing the bonus lab.

TABLE I: Lab Schedule

Date	Task
13th Nov	Lab 1: Lab Introduction & Programming Languages
20th Nov	Lab 2: Hardware Selection
27th Nov	Lab 3: Sparse Data
4th Dec	Lab 4: Frequency & Voltage
11th Dec	Extra Time + Bonus Lab

The labs work as follow: before each lab, you need to do a preliminary task, and reply the questions asked in the corresponding Canvas assignment. Please make sure to upload the answer to this preliminary task before coming to the lab, as the deadline for the assignment will be set at the time the corresponding session starts. By the start of the session, the TA will unlock in Canvas the resources required for the corresponding lab if needed. Once in the lab, follow the instructions given in this manual. During the lab session the TAs will ask you some questions, in a demo-like fashion. For that, for each lab a representative will be randomly selected among the members of the group (although we will make sure that everyone gets to be a representative). This representative will be the one in charge of responding to the TA's questions. After this questioning, please make sure that the TA has marked this demo as completed in Canvas. Finally, fill and submit the post-lab assignment in Canvas.

## I. LAB 1: LAB INTRODUCTION & PROGRAMMING LANGUAGES

General Matrix-Matrix Multiplication (GEMM) is one of the most extensively used kernels across multiple domains and applications. One such example are Convolution Neural Networks (CNNs), which can be used for tasks like image

TABLE II: Inference results for the provided images

#images	Statistic	Accuracy
1	Top 1	0%
	Top 5	100%
10	Top 1	20%
	Top 5	70%
100	Top 1	59%
	Top 5	85%

recognition. In this lab, you will focus on optimizing image recognition over the ImageNet dataset [1] using AlexNet [3], one of the most well known CNN models.

To run AlexNet, first get the pre-trained weights running `./get_weights.sh` (if it is not working at the beginning, try making sure you have execution permissions by running `chmod +x ./get_weights.sh`). After downloading the weights, make sure you have the right paths, as explained in the *README.md* file. Then you can run AlexNet with `python3 main.py`. When running this command, the program will do as many inferences (i.e., will recognize as many images) as set up in the constant `BATCH_SIZE`. We have provided 100 images in which you can infer in the *data* folder. The results for inferring these images with AlexNet are as shown in Table II.

To measure CPU power, we read the power and energy counters using Intel's RAPL. For that, IT has provided a binary that reads them and prints them to a file. To use this binary, run `sudo /chalmers/sw/sup64/phc/b/binh/rapl-read` in a terminal, and then launch AlexNet in a second terminal. When AlexNet finishes, go back to the terminal where power was being measured and kill the process with `Ctrl+C`. Then run `power_tools/energy_parser.py` with the start and end times you obtained from running AlexNet, and average power and total energy consumption will be printed.

### A. Pre-Lab Assignment

Read the paper [4] and estimate how much energy (in percentage) you could save by moving the computation of GEMM from Python to C++.

In addition to that, once you get to the lab, start running AlexNet with batch size 1 as soon as you can with `python3 main.py`. The baseline version takes quite a lot of time to run, so you can leave that on the background while reading the code and understanding the problem. To set up the batch size, change the corresponding value in *main.py*.

## B. Lab Assignment

As explained before, once you get to the lab try to get AlexNet running with batch size 1. The base version you were provided uses a manual implementation of GEMM developed in Python, and you will see the problems of it.

However, you do not need to run everything with Python. You offload parts to the more efficient C++ using Pybind11 [2]. In `src/cpp/gemm.cpp`, you have the template for doing so. Write your GEMM implementation in C++ there, and then build the library with `./compile_library.sh`. Then load this library in `gemm.py`, as use your C++ function instead of the Python one. How much energy are you saving with this change? Then, add parallelism in your C++ implementation with OpenMP. Are you saving energy with this new change? Why (not)? If you change the number of threads assigned, do you see any difference?

## REFERENCES

- [1] J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li, and L. Fei-Fei. Imagenet: A large-scale hierarchical image database. In *2009 IEEE conference on computer vision and pattern recognition*, pages 248–255. Ieee, 2009.
- [2] W. Jakob, J. Rhinelander, and D. Moldovan. pybind11 – seamless operability between c++11 and python, 2017. <https://github.com/pybind/pybind11>.
- [3] A. Krizhevsky, I. Sutskever, and G. E. Hinton. Imagenet classification with deep convolutional neural networks. *Advances in neural information processing systems*, 25, 2012.
- [4] R. Pereira, M. Couto, F. Ribeiro, R. Rua, J. Cunha, J. P. Fernandes, and J. Saraiva. Ranking programming languages by energy efficiency. *Science of Computer Programming*, 205:102609, 2021.